CaDoD: Phase 4 Cats vs Dogs Detector

Group 25

Martin Breth
mbreth@iu.edu

Vicente De Leon
vdeleonw@iu.edu

Kelly Craig
craigke@iu.edu

Courtney Payton
coupayto@iu.edu

CNN Model

- Dependencies: TensorFlow, Keras, Pandas, NumPy, Sklearn.
- Input dimensions: 128x128x3(Red,Green,Blue).
- ImageDataGenerator for data augmentation/preparation.
- ReLU, Softmax, and Binary Cross Entropy Loss.
- Adam Optimizer vs RMSprop Optimizer = Adam (winner).
- Images are based on Adam Optimizer.
- 25 epochs: 76.2% train acc, 75.9% validation acc, 76% test acc.

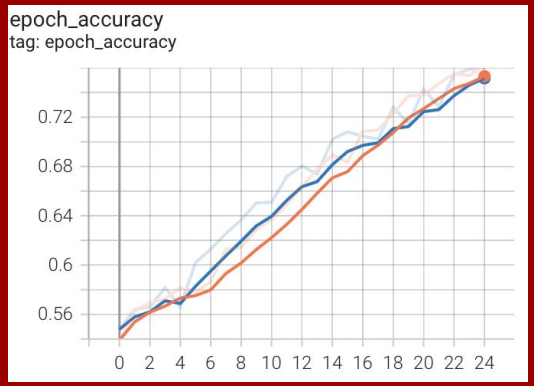

CNN Summary

```
Model: "sequential"
_________________________________________________________________
 Layer (type)                Output Shape              Param #
=================================================================
 conv2d (Conv2D)             (None, 128, 128, 32)      896

 conv2d_1 (Conv2D)           (None, 128, 128, 32)      9248

 max_pooling2d (MaxPooling2D  (None, 64, 64, 32)       0
 )

 conv2d_2 (Conv2D)           (None, 64, 64, 64)        18496

 conv2d_3 (Conv2D)           (None, 64, 64, 64)        36928

 max_pooling2d_1 (MaxPooling  (None, 32, 32, 64)       0
 2D)

 conv2d_4 (Conv2D)           (None, 32, 32, 128)       73856

 conv2d_5 (Conv2D)           (None, 32, 32, 128)       147584

 max_pooling2d_2 (MaxPooling  (None, 16, 16, 128)      0
 2D)

 conv2d_6 (Conv2D)           (None, 16, 16, 256)       295168

 conv2d_7 (Conv2D)           (None, 16, 16, 256)       590080

 max_pooling2d_3 (MaxPooling  (None, 8, 8, 256)        0
 2D)

 flatten (Flatten)           (None, 16384)             0

 dense (Dense)               (None, 256)               4194560

 dropout (Dropout)           (None, 256)               0

 dense_1 (Dense)             (None, 256)               65792

 dropout_1 (Dropout)         (None, 256)               0

 dense_2 (Dense)             (None, 2)                 514

=================================================================
Total params: 5,433,122
Trainable params: 5,433,122
Non-trainable params: 0
```

CNN Phase 4 vs MLP Phase 3

CNN Model (Adam Optimizer) Highest Accuracy Phase 4:

- Using 25 epochs.
- 76.2% train acc, 69% validation acc, 76% test acc.
- Major improvement since Phase 3.
- Actually classifies both cat and dog images.
- Section 4.5.5 of the notebook compares both Phases.
- It compares classification report and conf. matrix.

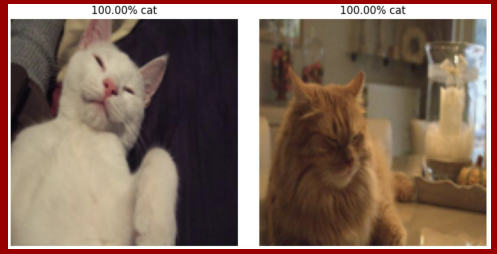

	precision	recall	f1-score	support
cat	0.82	0.29	0.43	1209
dog	0.60	0.94	0.74	1385
accuracy			0.64	2594
macro avg	0.71	0.62	0.58	2594
weighted avg	0.70	0.64	0.59	2594

MLP Model (Adam Optimizer) Highest Accuracy Phase 3:

- Using 20 epochs.
- 55.14% train acc, 56.62% validation acc, 55.28% test acc.
- Having more trouble classifying almost all cat images.
- See section 4.5.5 of the notebook.

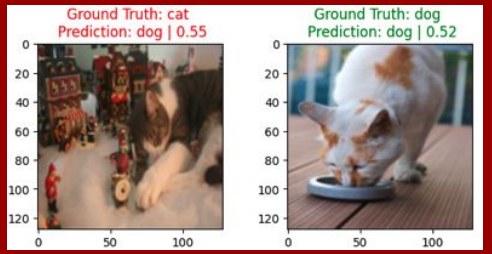

	precision	recall	f1-score	support
cat	0.71	0.07	0.13	1209
dog	0.55	0.98	0.70	1385
accuracy			0.55	2594
macro avg	0.63	0.52	0.41	2594
weighted avg	0.62	0.55	0.43	2594

Results CNN

CNN Adam Optimizer highest scores:

- 25 epochs.
- We can clearly see Phase 4 CNN has higher scores.

	Adam Train Accuracy	Adam Train loss	Adam Validation Accuracy	Adam Validation Loss
0	0.762	0.69	0.76	0.69

```
The test loss for CNN Model is: 0.487
The test accuracy for CNN Model is: 0.76
```

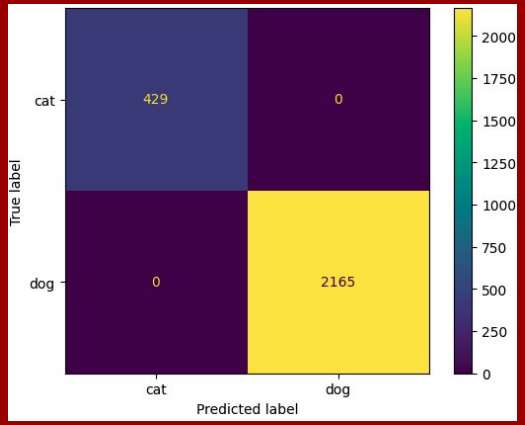

MLP Adam Optimizer highest scores:

- 20 epochs.
- Even though MLP had a 55% test accuracy, the model wasn't performing correctly.
- Possibly the most inaccurate model of all classifiers.

	MLP Optimizer	Number of Epoch	Train Accuracy	Validation Accuracy	Test Accuracy
Adam Optimizer - cpu		epoch=20	55.140	56.627	55.281

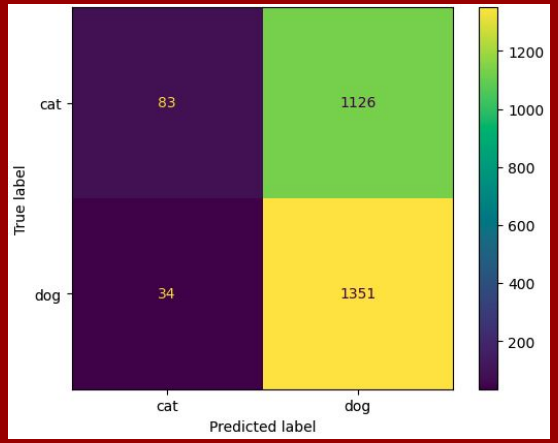

Fully Convolutional Networks (FCN)

- Similar to CNN
- Comparing Adam and SGD optimizers with different batch sizes
- More epochs needed

- TensorBoard was used for real-time results

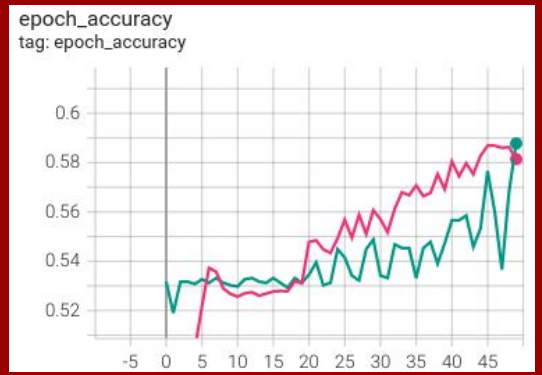

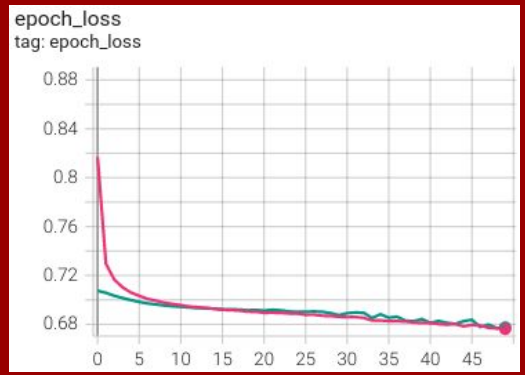

	Model	Epoch	Train Accuracy	Valid Accuracy	Test Accuracy	Train Loss	Valid Loss	Test Loss
0	FCN (Adam, batch 25)	46	0.622	0.557	0.592	0.671	0.681	0.676
1	FCN (SGD, batch 50)	48	0.516	0.528	0.529	0.776	0.713	0.712
2	FCN (Adam, batch 50)	45.0	0.587	0.577	0.592	0.68	0.684	0.676

AlexNet

- Reshaped data input size to (227,227,3)
- Rotated images and applied a flip to avoid overfitting
- Fine-tuned model through testing different learning rates and optimizers(SGD, and Adam)

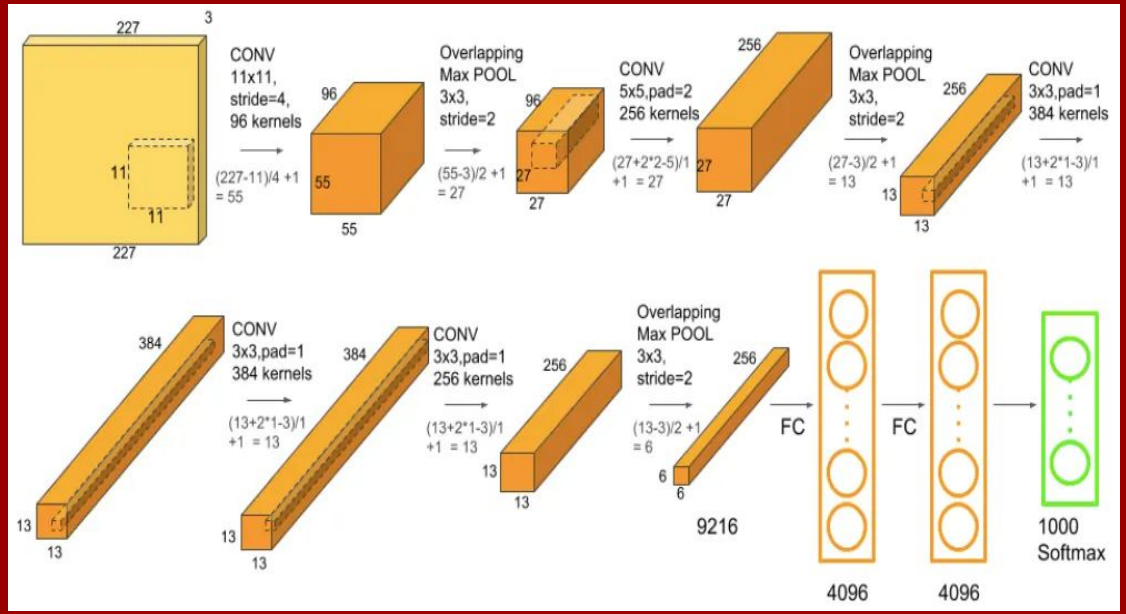

AlexNet Results - Best Model with SGD and Learning rate of 0.1

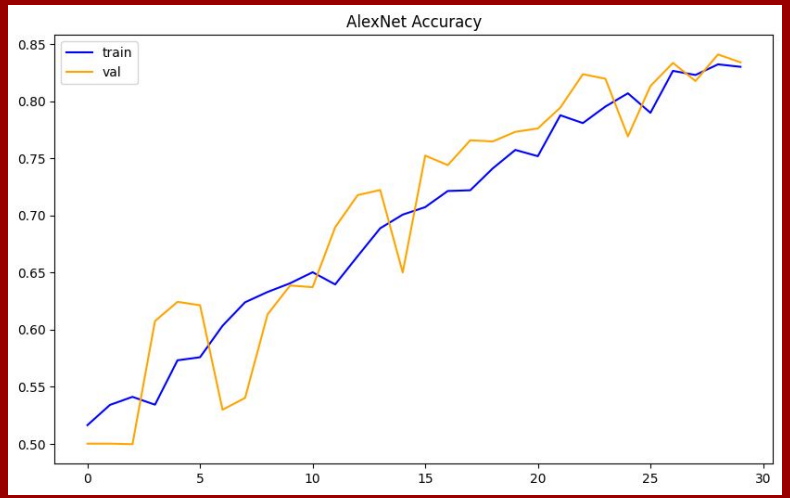

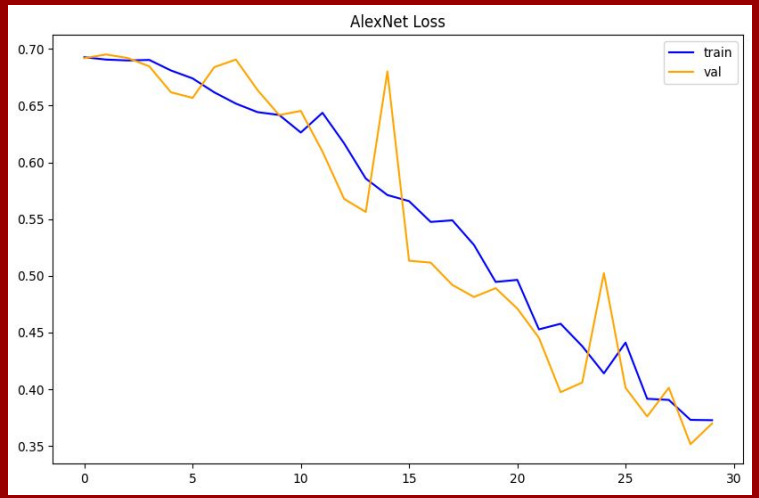

30 epochs	Accuracy	Loss
Train	81.57%	0.3866
Validation	84.73%	0.3413

Transfer Learning with EfficientNet

- Image Classification Family of Models
- EfficientNet is pre-trained on ImageNet data.
- Froze EfficientNet Model and added a top layer
- Experiment with Regularization
- Optimizer = RMSprop
- Loss function = Binary Cross Entropy

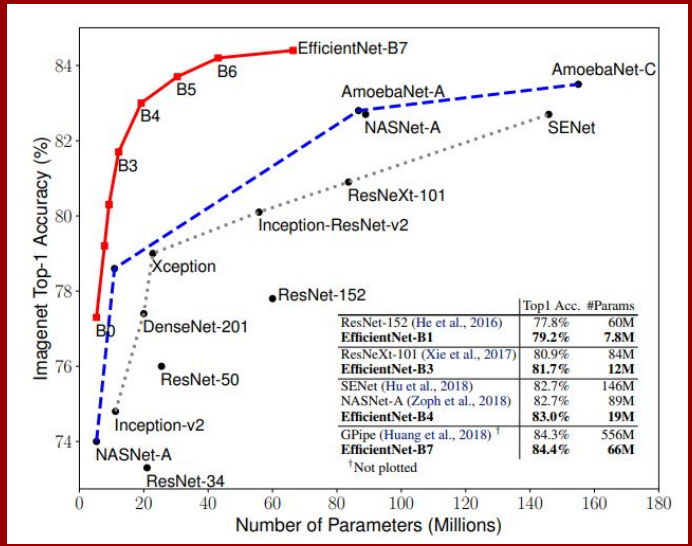

EfficientNet

- Best Model:
  - EfficientNet B5
  - Regularization - weight decay of 1e-6.
- Problems:
  - Memory limitations for EfficientNet B6 and EfficientNet B7

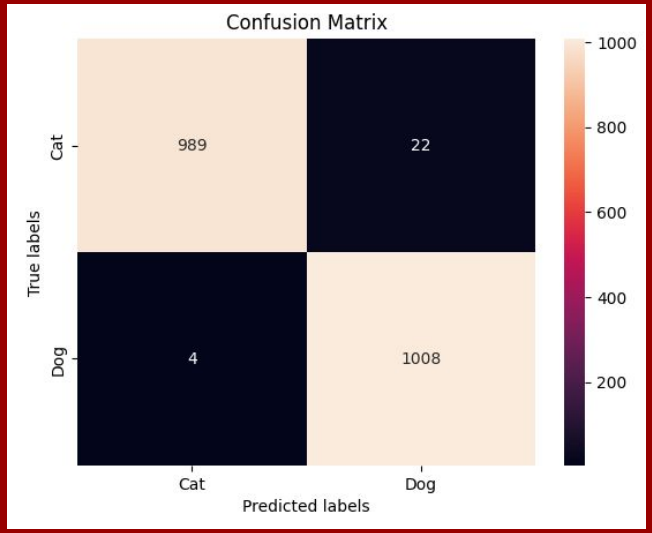

	Base Model	Input Resolution	Optimizer	Epochs	Training Accuracy	Validation Accuracy	Test Accuracy	Fit Time	Total Parameters
0	EfficientNet B0	224	RMSprop	10	0.951500	0.98750	0.971330	2195.619411	68276893
1	EfficientNet B0 with decay	224	RMSprop	10	0.956171	0.98750	0.974790	381.655278	68276893
2	EfficientNet B1	240	RMSprop	10	0.959698	0.99125	0.985171	349.486195	90463361
3	EfficientNet B2	260	RMSprop	10	0.965239	0.99500	0.983193	380.205944	124555763
4	EfficientNet B3	300	RMSprop	10	0.970277	0.99750	0.988631	516.328792	168071977
5	EfficientNet B4	380	RMSprop	10	0.973300	0.99875	0.987148	787.670477	281917017
6	EfficientNet B5	456	RMSprop	10	0.979849	0.99875	0.994563	1203.512686	500374769

Results and Discussion

CNN:

- 76.2% train acc, 69% validation acc, 76% test acc.
- Major improvement since Phase 3.

FCN:

- Best model: Adam Optimizer gave best results
- Accuracy: 58.7% train 57.7% valid and 59.2% test

AlexNet:

- Best Model: Stochastic Gradient Descent (SGD), and learning rate: 0.1
- 81.6% train acc, 84.7% validation acc
- Major improvement since Phase 3

EfficientNet

- Best Model: Rmsprop optimizer, EfficientNetB5
- 97.9% train acc, 99.9 val acc, and 99.5 test accuracy